

MATHEMATICAL STATISTICS

1. BASIC SETUP

Definition 1.1 (Statistical setup). A *statistical setup* consists of the following data:

We have the following measurable spaces:

- The *observation space* $\mathcal{X}$.
- The *parameter space* Θ .
- The *action space* $\mathcal{A}$.
- The *sample space* Ω

We also have the following maps:

- A *loss function* $L : \Theta \times \mathcal{A} \rightarrow \mathbb{R}$.
- An *observation* $X : \Omega \rightarrow \mathcal{X}$.
- A *decision function* $d : \mathcal{X} \rightarrow \mathcal{A}$.

For each parameter $\theta \in \Theta$, we have a probability measure $P_\theta \in \text{Prob}(\Omega)$. The associated risk function is then

$$\begin{aligned} R(\theta, d) &= E_\theta L(\theta, d(X)) \\ &= \int_{\Omega} L(\theta, d(X(\omega))) dP_\theta(\omega). \end{aligned}$$

Note: It seems that unless otherwise specified, it does not hurt to just assume $\mathcal{X} = \Omega$ and X is the identity map in which case

$$R(\theta, d) = \int_{\mathcal{X}} L(\theta, d(x)) dP_\theta(x).$$

Interpretation: We have an unknown parameter θ and a given observation $X(\omega) \in \mathcal{X}$ where $\omega \in \Omega$ was randomly sampled according to P_θ . Based on this observation, we take an action $d(x) \in \mathcal{A}$. Then the loss $L(\theta, d(x))$ measures how wrong we are in taking this action $d(x)$. The risk $R(\theta, d)$ is then the average of this loss (for the fixed θ).

Example 1.2. Suppose now $\Omega = \mathcal{X} = \mathbb{R}^n$ and X is identity (as usual) and $\mathcal{A} = \mathbb{R}$. Let $\Theta = \mathbb{R} \times \mathbb{R}_{\geq 0}$ and for $\theta = (\mu, \sigma^2) \in \Theta$ we let P_θ denote the distribution of the random variable $(X_1, \dots, X_n) \in \mathbb{R}^n$ where the X_i are i.i.d $N(\mu, \sigma^2)$. We take a loss function $L(\theta, a) = |\mu - a|^2$. Thus if $d : \mathcal{X} \rightarrow \mathcal{A}$ is a decision rule then $L(\theta, d(x))$ will measure how well of an approximation $d(x) \in \mathbb{R}$ is of the unknown μ . Let us choose $d(x_1, \dots, x_n) = \frac{1}{n} \sum_{i=1}^n x_i$. Thus

$$R(\theta, d) = E \left| \frac{1}{n} \sum_{i=1}^n X_i - \mu \right|^2 = \text{Var} \left(\frac{1}{n} (X_1 + \dots + X_n) \right) = \frac{\sigma^2}{n}$$

by the independence of the X_i . We see that as $R(\theta, d) = O(n^{-1})$ so this seems to be a good decision rule.

If instead we use the loss function $L(\theta, a) = |\sigma^2 - a|$ then a good decision rule seems to be

$$d(X) = s^2 := \frac{1}{n-1} \sum_{i=1}^n |X_i - \bar{X}|^2.$$

It turns out that (TODO?)

$$R(\theta, d) = \frac{2\sigma^4}{n-1}.$$

We can use this language to define hypothesis testing:

Definition 1.3. Suppose now that we have a partition $\Theta = \Theta_0 \sqcup \Theta_1$ of the parameter space and our action space is $\mathcal{A} = \{\Theta_0, \Theta_1\}$. We have a 0 – 1 loss function given by

$$L(\theta, \Theta_i) = 1 - \mathbf{1}_{\Theta_i}(\theta).$$

Thus $L(\theta, a)$ measures whether we chose the right partition $a \in \mathcal{A}$ that θ landed in. We see that the risk

$$R(\theta, d) = P_\theta(\{\omega \in \Omega \mid \theta \notin d(X(\omega))\})$$

is the probability that our decision algorithm chooses the wrong partition.

The action Θ_0 is called the *Null Hypothesis* while Θ_1 is called the *Alternative hypothesis*. A *Type I* error is an action that chooses the alternative hypothesis (Θ_1) when the null hypothesis is true ($\theta \in \Theta_0$). A *Type II error* is an action that chooses the null hypothesis when the alternative is true.

Example 1.4. Suppose we have a dice that is either fair (null hypothesis) or always rolls a 6 (alternative hypothesis). Suppose we have rolled this dice n times and observed the outcomes $(X_1, \dots, X_n) \in \{1, 2, 3, 4, 5, 6\}^n = \mathcal{X}$. Say $\Theta = \{\theta_0, \theta_1\}$ where P_{θ_0} is the uniform distribution on $\mathcal{X}$ (dice is fair) while P_{θ_1} is the probability measure concentrated on $(6, \dots, 6) \in \mathcal{X}$. Our decision rule simply chooses θ_1 if and only if we roll all 6. We have $R(\theta_0, d) = 1/6^n$ while $R(\theta_1, d) = 0$.

Definition 1.5. We say that a statistical setup is *nice* if there exists a σ -finite positive measure (not necessarily a probability measure) ν on $\Omega = \mathcal{X}$ and a measurable function $f : \mathcal{X} \times \Theta \rightarrow \mathbb{R}_{\geq 0}$ such that $P_\theta \ll \nu$ and

$$\frac{dP_\theta}{d\nu}(x) = f(x|\theta) := f(x, \theta)$$

for all $\theta \in \Theta$.

For example, if $\mathcal{X}$ is countable and thus each P_θ is discrete, then ν can be the counting measure.

Assumption going forward: We always assume a statistical setup is nice without mentioning it. This is the case for the examples above (i.e., ν is either a Lebesgue measure or a counting measure).

2. BAYESIAN SETUP

Definition 2.1. A *prior distribution* is a probability measure μ_Θ on the parameter space Θ .

Recalling that we assume the nice condition of Definition 1.5, we thus have a probability measure P on $\mathcal{X} \times \Theta$ given by

$$\frac{dP}{d(\nu \times \mu_\Theta)}(x, \theta) = f(x|\theta).$$

In other words for integrable $\psi : \mathcal{X} \times \Theta \rightarrow \mathbb{R}$ we have

$$\begin{aligned}
\int_{\mathcal{X} \times \Theta} \psi(x, \theta) dP(x, \theta) &= \int_{\mathcal{X} \times \Theta} \psi(x, \theta) f(x|\theta) d(\nu \times \mu_\Theta)(x, \theta) \\
&= \int_{\Theta} \left(\int_{\mathcal{X}} \psi(x, \theta) f(x|\theta) d\nu(x) \right) d\mu_\Theta(\theta) \\
&= \int_{\Theta} \int_{\mathcal{X}} \psi(x, \theta) dP_\theta(x) d\mu_\Theta(\theta)
\end{aligned}$$

From now on, we let $X : \mathcal{X} \times \Theta \rightarrow \mathcal{X}$ and $\Theta : \mathcal{X} \times \Theta \rightarrow \Theta$ denote the projections (note: by abuse of notation Θ is now either a space or a projection map, but it should cause no confusion in any context).

Let $P_X = X^*P$ denote the push-forward measure of P onto $\mathcal{X}$, that is, $P_X(A) = P(A \times \Theta)$ for $A \subset \mathcal{X}$. We let

$$g(x) := \int_{\Theta} f(x, \theta) d\mu_\Theta(\theta).$$

Observe that for $A \subset \mathcal{X}$ we have that

$$P_X(A) = \int_{A \times \Theta} f(x|\theta) d\mu_\Theta(\theta) d\nu(x) = \int_A g(x) d\nu(x)$$

and thus P_X has a density with respect to ν given by g .

Lemma 2.2. For P_X -almost all $x \in \mathcal{X}$ we have that

$$g(x) > 0.$$

Proof. Let $A = \{x \in \mathcal{X} \mid g(x) = 0\}$. Then

$$P_X(A) = \int_A g(x) d\nu(x) = 0.$$

□

Given this lemma, we can now define for P_X -almost all $x \in \mathcal{X}$ a probability measure on Θ given by

$$P_{\Theta|X=x}(B) = \frac{1}{g(x)} \int_B f(x|\theta) d\mu_\Theta(\theta)$$

for $B \subset \Theta$. This is known as the *posterior distribution of Θ given $X = x$* .

Proposition 2.3. The measures $\delta_x \times P_{\Theta|X=x}$ are conditional distributions of P given $X = x$ where $X : \mathcal{X} \times \Theta \rightarrow \mathcal{X}$ is the projection. The conditional expectation of an integrable $\psi : \mathcal{X} \times \Theta \rightarrow \mathbb{R}$ given $X = x$ is given by

$$E(\psi|X = x) = \frac{1}{g(x)} \int_{\Theta} \psi(x, \theta) f(x|\theta) d\mu_\Theta(\theta).$$

Proof. Let $E'(\psi|X = x)$ denote the right hand side. We first show that it is well defined for almost all $x \in \mathcal{X}$. This follows since by assumption we have that $|\psi(x, \theta)|$ is integrable with respect to P and thus

$$\int |\psi(x, \theta)| f(x|\theta) d\nu(x) d\mu_\Theta(\theta) < \infty.$$

It follows from Fubini-Tonelli that for ν almost all $x \in \mathcal{X}$ we have that

$$\int_{\mathcal{X}} |\psi(x, \theta)| f(x|\theta) d\mu_\Theta(\theta) < \infty$$

and thus $E'(\psi|X = x)$ is a well defined measurable function on X . We have to show that $(x, \theta) \mapsto E'(\psi|X = x)$ is measurable with respect to the σ -algebra induced by $X : \mathcal{X} \times \Theta \rightarrow \mathcal{X}$ and that for all $A \subset X$ we have that

$$\int_{A \times \Theta} E'(\psi|X = x) dP(x, \theta) = \int_{A \times \Theta} \psi(x, \theta) dP(x, \theta).$$

The measurability is obvious as there is no dependence on Θ . For the second claim, note that since (as shown above) $g = dP_X/d\nu$, we have

$$\begin{aligned} \int_{A \times \Theta} E'(\psi|X = x) dP(x, \theta) &= \int_A E'(\psi|X = x) dP_X(x) \\ &= \int_A E'(\psi|X = x) g(x) d\nu(x) \\ &= \int_A \int_{\Theta} \psi(x, \theta) f(x|\theta) d\mu_{\Theta}(\theta) d\nu(x) \end{aligned}$$

which completes the proof of this property since $f(x|\theta) = \frac{dP}{d(\nu \times \mu_{\Theta})}(x, \theta)$. Finally, observe that

$$\int_{\mathcal{X} \times \Theta} \psi d(\delta_x \times P_{\Theta|X=x}) = \int_{\Theta} \psi(x, \theta) dP_{\Theta|X=x}(\theta) = \frac{1}{g(x)} \int_{\Theta} f(x|\theta) \psi(x, \theta) d\mu_{\Theta}(\theta) = E(\psi | X = x)$$

so indeed $\delta_x \times P_{\Theta|X=x}$ is the conditional probability of P given $X = x$. $\square$

Proposition 2.4. For $A \subset \mathcal{X}$ we have that

$$P(A \times \Theta | \Theta = \theta) = P_{\theta}(A).$$

Proof. By symmetry, if we swap the roles of $\mathcal{X}$ with Θ and apply the proposition above, then we get

$$\begin{aligned} P(A \times \Theta | \Theta = \theta) &= \left(\int_{\mathcal{X}} f(x|\theta) d\nu(x) \right)^{-1} \left(\int_A f(x|\theta) d\nu(x) \right) \\ &= P_{\theta}(X)^{-1} P_{\theta}(A) \\ &= P_{\theta}(A). \end{aligned}$$

$\square$

2.1. Notation. We use the standard notation from probability theory: if $F_i : \Omega \rightarrow \Omega_i$ are measurable and P is a measure on Ω , then

$$P(F_1 \in A_1, \dots, F_n \in A_n) := P(\{\omega \in \Omega \mid F_1(\omega) \in A_1, \dots, F_n(\omega) \in A_n\})$$

and likewise for conditional measures.

Thus if P is the measure on $\mathcal{X} \times \Omega$ as constructed above, then

$$P(X \in A, \Theta \in B) = P(A \times B).$$

3. SUFFICIENT STATISTICS

3.1. Basic definitions and examples. A statistic is a measurable function $T : \mathcal{X} \rightarrow \mathcal{T}$ to some measurable space $(\mathcal{T}, \mathcal{B}_{\mathcal{T}})$.

Definition 3.1. We say that a statistic T on $\mathcal{X}$ is *sufficient (in the Bayesian sense)* if for any prior μ_{Θ} on Θ we have for any Borel $B \subset \Theta$ we have the equality

$$P(\Theta \in B \mid X = x) = P(\Theta \in B \mid T \circ X = Tx) \quad \text{for } \mu_X\text{-almost all } x \in \mathcal{X}.$$

Example 3.2. Consider the case where $\mathcal{X} = \{0, 1\}^n$ and $\Theta = [0, 1]$. Suppose that P_θ is the Binomial probability distribution on $\{0, 1\}^n$ where $(X_1, \dots, X_N) \in \mathcal{X}$ is drawn i.i.d with $P(X_i = 1) = \theta$. Let us for now assume the Prior distribution μ_Θ is unknown. We know that

$$P_\theta(X = x) = \theta^{T(x)}(1 - \theta)^{n-T(x)}$$

where $T(x) = \sum_{i=1}^n x_i$ is the number of successes. The posterior distribution is thus given by

$$P(\Theta \in B \mid X = x) = \frac{P(\Theta \in B, X = x)}{P(X = x)} = \frac{\int_B \theta^{T(x)}(1 - \theta)^{n-T(x)} d\mu_\Theta(\theta)}{\int_\Theta \theta^{T(x)}(1 - \theta)^{n-T(x)} d\mu_\Theta(\theta)}$$

for all $x \in \mathcal{X}$ that occur with positive probability (this is all of $\mathcal{X}$ except for possibly $x = (0, \dots, 0)$ or $x = (1, \dots, 1)$ in case θ is concentrated on 1 or 0 respectively).

We claim that $T : \mathcal{X} \rightarrow \mathbb{Z}_{\geq 0}$ given by $T(x) = \sum_{i=1}^n x_i$ is a sufficient statistic. Indeed if we fix $x_0 \in \mathcal{X} = \{0, 1\}^n$ and we let $m = Tx_0$, we get that

$$\begin{aligned} P(\Theta \in B \mid T \circ X = m) &= \frac{P(\Theta \in B, T \circ X = m)}{P(T \circ X = m)} \\ &= \frac{\sum_{Tx=m} P(\Theta \in B, X = x)}{\sum_{Tx=m} P(X = x)} \\ &= \frac{\sum_{Tx=m} \int_B \theta^{T(x)}(1 - \theta)^{n-T(x)} d\mu_\Theta(\theta)}{\sum_{Tx=m} \int_\Theta \theta^{T(x)}(1 - \theta)^{n-T(x)} d\mu_\Theta(\theta)} \\ &= \frac{\binom{n}{m} \int_B \theta^m (1 - \theta)^{n-m} d\mu_\Theta(\theta)}{\binom{n}{m} \int_\Theta \theta^m (1 - \theta)^{n-m} d\mu_\Theta(\theta)} \end{aligned}$$

And so we see that this fraction cancels out to the expression for $P(\Theta \in B \mid X = x_0)$ that we already obtained.

In fact, the same argument shows that T is still a sufficient statistic if instead of assuming that P_θ is binomial we assume that P_θ is *exchangeable*, meaning a distribution on $\mathcal{X} = \{0, 1\}^n$ that is invariant under permutations of the co-ordinates, in this case the same calculation will instead yield

$$P(\Theta \in B \mid TX = m) = \frac{\sum_{Tx=m} \int_B P_\theta(x) d\mu_\Theta}{\sum_{Tx=m} \int_\Theta P_\theta(x) d\mu_\Theta}$$

but since we assume that $P_\theta(x) = P_\theta(x')$ whenever x and x' are permutations of each other, we must have that all the summands are the same and hence we get the same cancellation.

Remark 3.3. Intuitively, sufficiency of a statistic essentially means that if we want to estimate the Posterior distribution of θ given an observable statistic, we only need to look at the statistic of $T \circ X$ rather than X , which may be simpler. In other words, somebody who wishes to estimate θ (e.g., find a confidence interval) need not know the sample x but just needs to know Tx . In this example, the person wishing to estimate θ needs only to be told a single integer Tx , rather than a potentially very long sequence $x \in \{0, 1\}^n$ (data compression).

The classical definition of sufficiency does not make reference to a prior distribution and is thus not Bayesian.

Definition 3.4. We say that a statistic $T : \mathcal{X} \rightarrow \mathcal{T}$ is *sufficient in the classical sense* if we have a function $r : \mathcal{B} \times \mathcal{T} \rightarrow [0, 1]$ such that for all $\theta \in \Theta$ and $B \in \mathcal{B}$ (recall $\mathcal{B}$ is the Borel σ -algebra on $\mathcal{X}$) we have that

$$P_\theta(B \mid T = t) = r(B, t) \quad \text{for } T^*P_\theta - \text{almost all } t \in \mathcal{T}$$

and $t \mapsto r(B, t)$ is Borel on $\mathcal{T}$.

This is a technical way of saying that T is sufficient if and only if $P_\theta(B|T = t)$ does not depend on θ . We stated the definition formally in order to deal with the issue that conditional probability is only almost everywhere unique and well defined.

Example 3.5. Let us continue from the binomial distribution in Example 3.2. As $\mathcal{X} = \{0, 1\}^n$ is a finite set, we only have to consider the case where $B \subset \mathcal{X}$ is a singleton $B = \{x_0\}$ for some $x_0 \in \mathcal{X}$. Now $P_\theta(B|T = t) = 0$ if $Tx_0 \neq t$. If $Tx_0 = t$ then we have that

$$P_\theta(B|T = t) = \frac{P_\theta\{x_0\}}{P_\theta\{x \in \mathcal{X} \mid Tx = x_0\}} = \frac{1}{\binom{n}{t}}$$

where the last equality follows from the fact that all $x \in \mathcal{X}$ with $Tx = t$ have the same probability in a binomial distribution. In fact, the same argument applies to the case where the P_θ are all exchangeable.

Thus we may take

$$r(\{x_0\}, t) = \mathbb{1}(Tx_0 = t) \binom{n}{t}^{-1}$$

and extend r to $\mathcal{B} \times \mathbb{Z}$ by additivity. Thus this T is indeed sufficient in the classical sense.

3.2. Equivalence of definitions. In this section we prove equivalence of the classical and Bayesian definition of a sufficient statistics under the nice assumptions of Definition 1.5 (which we always assume).

Lemma 3.6. Suppose $T : \mathcal{X} \rightarrow \mathcal{T}$ is a statistic and μ_Θ is a prior on Θ . Suppose also that there is a Borel map $s : \mathcal{T} \times \Theta \rightarrow [0, 1]$ so that

$$s(t, \theta) = P_\theta(A \mid T = t) \quad \text{for } T^*P_\theta\text{-almost all } t \in \mathcal{T}.$$

We have that

$$P(X \in A \mid TX = Tx, \Theta = \theta) = s(Tx, \theta) \quad \text{for } P\text{-almost all } (x, \theta) \in \mathcal{X} \times \Theta.$$

Proof. We must show that for $B \subset \Theta, C \subset \mathcal{T}$ we have that

$$\int_{T^{-1}C \times B} P(A \times \Theta \mid Tx = t, \Theta = \theta) dP(x, \theta) = \int_{T^{-1}C \times B} s(Tx, \theta) dP(x, \theta).$$

The law of total probability tells us that the left hand side is equal to

$$P((A \times \Theta) \cap (T^{-1}C \times B)) = P(X \in A \cap T^{-1}C, \Theta \in B).$$

The right hand side is

$$\begin{aligned} \int_B \left(\int_{T^{-1}C} P_\theta(A|T = Tx) dP_\theta(x) \right) d\mu_\Theta(\theta) &= \int_B P_\theta(A \cap T^{-1}C) d\mu_\Theta(\theta) \\ &= \int_B \int_{\mathcal{X}} \mathbb{1}_{A \cap T^{-1}C}(x) dP_\theta(x) d\mu_\Theta(\theta) \\ &= P(X \in A \cap T^{-1}C, \Theta \in B) \end{aligned}$$

□

Theorem 3.7. Suppose that $T : \mathcal{X} \rightarrow \mathcal{T}$ is sufficient in the classical sense. Then T is also sufficient in the Bayesian sense (for any prior μ_Θ on Θ). Moreover for $A \subset \mathcal{X}$ we have

$$(1) \quad P(X \in A \mid TX = Tx, \Theta = \theta) = P(X \in A \mid TX = Tx) \quad \text{for } P\text{-almost all } (x, \theta) \in \mathcal{X} \times \Theta.$$

Proof. We first show the *Moreover* claim. By the definition of sufficiency we have that

$$P_\theta(A \mid T = t) = r(A, t)$$

for some measurable $t \mapsto r(A, t)$ and so indeed the Borel condition of the previous Lemma is satisfied. By the previous Lemma the left hand side of (1) is equal to $r(A, Tx)$. So now it remains to show that $r(A, Tx) = P(A \times \Theta \mid TX = Tx)$ for P -almost all (x, θ) . Thus we must show that for $C \subset \mathcal{T}$ we have that

$$(2) \quad \int_{T^{-1}C \times \Theta} r(A, Tx) dP(x, \theta) = \int_{T^{-1}C \times \Theta} P(A \times \Theta \mid TX = Tx) dP(x, \theta).$$

By definition of conditional probability w.r.t $TX : \mathcal{X} \times \Theta \rightarrow \mathcal{T}$, the right hand side is equal to

$$P(A \times \Theta \cap T^{-1}C \times \Theta) = P(X \in A \cap T^{-1}C).$$

Repeating the calculation of the previous Lemma we have that the left hand side of (2) is

$$\begin{aligned} \int_{\Theta} \left(\int_{T^{-1}C} P_\theta(A \mid T = Tx) dP_\theta(x) \right) d\mu_\Theta(\theta) &= \int_{\Theta} P_\theta(A \cap T^{-1}C) d\mu_\Theta(\theta) \\ &= \int_{\Theta} \int_{\mathcal{X}} \mathbb{1}_{A \cap T^{-1}C}(x) dP_\theta(x) d\mu_\Theta(\theta) \\ &= P(X \in A \cap T^{-1}C). \end{aligned}$$

This completes the proof of (2). Now it remains to show that T is sufficient in the Bayesian sense. By this just proved *Moreover* claim, it follows from Lemma D.4 that X and Θ are conditionally independent given TX . Using Lemma D.4 again we have that

$$P(\Theta \in B \mid TX = Tx, X = x) = P(\Theta \in B \mid TX = Tx) \quad \text{for } P\text{-almost all } (x, \theta)$$

and thus for P_X -almost all $x \in \mathcal{X}$. But the left hand side is clearly $P(\Theta \in B \mid X = x)$. Thus T is indeed sufficient in the Bayesian sense. $\square$

Lemma 3.8. A statistic $T : \mathcal{X} \rightarrow \mathcal{T}$ is sufficient in the Bayesian sense if and only if for all prior distributions μ_Θ on Θ and measurable $B \subset \mathcal{X}$ we have that

$$x \mapsto P(\Theta \in B \mid x \in X)$$

is almost surely $T^{-1}\mathcal{B}_T$ -measurable (i.e., almost surely equal to a function of T).

Proof. We prove the non-obvious direction only: Thus suppose that for each measurable $B \subset \mathcal{X}$ we have that $x \mapsto P(\Theta \in B \mid x \in X) = h_B(Tx)$ where $h_B : \mathcal{T} \rightarrow \mathbb{R}$ is Borel. We first use apply the law of total probability to the map $T \circ X : \mathcal{X} \times \Theta \rightarrow \mathcal{T}$ as follows. For $A \subset \mathcal{T}$ Borel we have that

$$P(\Theta \in B, T \circ X \in A) = \int_A P(\Theta \in B \mid T \circ X = t) dP_T(t).$$

Now we apply the law of total probability to the map $X : \mathcal{X} \times \Theta \rightarrow \mathcal{X}$ to obtain that

$$\begin{aligned}
P(\Theta \in B, T \circ X \in A) &= P(\Theta \in B, X \in T^{-1}A) \\
&= \int_{T^{-1}A} P(\Theta \in B | X = x) dP_X(x) \\
&= \int_{T^{-1}A} h_B(Tx) dP_X(x) \\
&= \int_A h_B(t) dP_{\mathcal{T}}(t).
\end{aligned}$$

Thus we must have that $P(\Theta \in B | T \circ X = t) = h_B(t)$ holds for almost all t . Thus

$$P(\Theta \in B | T \circ X = Tx) = h_B(Tx) = P(\Theta \in B | X = x)$$

holds for almost all $x \in X$. $\square$

Lemma 3.9. If $T : \mathcal{X} \rightarrow \mathcal{T}$ is sufficient in the Bayesian sense then there exists a probability measure ν^* on $\mathcal{X}$ such that $P_\theta \ll \nu^*$ for all $\theta \in \Theta$, $\nu^* \ll \nu$ and for each $\theta \in \Theta$ there is a Borel

$$h_\theta : \mathcal{T} \rightarrow [0, \infty)$$

such that

$$\frac{dP_\theta}{d\nu^*}(x) = h_\theta(Tx) \quad \text{for } \nu^*\text{-almost all } x.$$

Proof. By Lemma E.2 there exists a countable or finite sequence of distinct $\theta_1, \theta_2, \dots$ and $0 < c_1, c_2, \dots \leq 1$ with $\sum_i c_i = 1$ such that $\nu^* := \sum_{i>0} c_i P_{\theta_i}$ satisfies that $P_\theta \ll \nu^* \ll \nu$ for all $\theta \in \Theta$. Now fix any $\theta_0 \in \Theta$. First suppose $\theta_0 \neq \theta_i$ for all $i > 0$. Define now a prior μ on Θ where $\mu(\{\theta_0\}) = 1/2$ and $\mu(\{\theta_i\}) = c_i/2$, i.e., μ is discrete. Now we have

$$\begin{aligned}
P(\Theta = \theta_0 | X = x) &= \frac{\int_{\{\theta_0\}} f(x|\theta) d\mu(\theta)}{\int_{\Theta} f(x|\theta) d\mu(\theta)} \\
&= \frac{\frac{1}{2} f(x|\theta_0)}{\frac{1}{2} f(x|\theta_0) + \frac{1}{2} \sum_{i>0} c_i f(x|\theta_i)}.
\end{aligned}$$

Observe that the pushforward measure on $\mathcal{X}$ is $P_X = \frac{1}{2} P_{\theta_0} + \frac{1}{2} \nu^*$. We know by Lemma 2.2 that for P_X -almost all $x \in X$ the denominator is non-zero. Now observe that

$$\frac{d\nu^*}{d\nu}(x) = \sum_{i>0} c_i f(x|\theta_i)$$

and as $\nu^* \ll \nu$ we have that this is ν^* almost surely non-zero and thus P_X -almost surely non-zero as $P_X \ll \nu^*$ as shown above. So we can divide by this quantity to get that

$$(3) \quad \frac{dP_{\theta_0}}{d\nu^*(x)} = f(x|\theta_0) \left(\sum_{i>0} c_i f(x|\theta_i) \right)^{-1} \quad \text{for } P_X\text{-almost all } x \in \mathcal{X}.$$

Thus we have that

$$P(\Theta = \theta_0 | X = x) = \frac{dP_{\theta_0}/d\nu^*(x)}{dP_{\theta_0}/d\nu^*(x) + 1}$$

holds P_X -almost surely. However, P_X -almost surely we have that the right hand side is $\mathcal{T}$ -measurable. Thus ν^* -almost surely it also is. This completes the proof in the case $\theta_0 \neq \theta_i$ for all $i > 0$. If $\theta_0 = \theta_i$ for some $i > 0$ we instead define the prior

$$\mu = \sum_{i>0} c_i P_{\theta_i}.$$

Without loss of generality, say $\theta_0 = \theta_1$. Thus

$$P(\Theta = \theta_1 \mid X = x) = \frac{c_1 f(x|\theta_1)}{\sum_{i>0} c_i f(x|\theta_i)} = c_1 \frac{dP_{\theta_1}}{d\nu^*}(x).$$

Thus again we have that this is P_X -almost equal to a $\mathcal{T}$ -measurable function (and thus ν^* almost equal to such).

□

Theorem 3.10. Assume the nice condition of Definition 1.5 . If $T : \mathcal{X} \rightarrow \mathcal{T}$ is sufficient in the Bayesian sense, then it is also sufficient in the classical sense. Letting ν^* be the measure from the previous Lemma, we have the more precise formulation of classical sufficiency: There is a $\mathcal{T}$ -measurable set $\mathcal{X}' \subset \mathcal{X}$ such that for all $\theta \in \Theta$ we have $\nu^*(\mathcal{X}') = 1 = P_\theta(\mathcal{X}')$ and for $A \subset X$ we have that

$$P_\theta(A \mid T = Tx) = \nu^*(A \mid T = Tx) \quad \text{for all } x \in \mathcal{X}'.$$

Proof. Let ν^* be the measure from the previous Lemma and recall that

$$\frac{dP_\theta}{d\nu^*}(x) = h_\theta(Tx).$$

We will now show that for all $\theta \in \Theta$ and $A \subset X$ we have that

$$P_\theta(A \mid T = t) = \nu^*(A \mid T = t),$$

which will complete the proof. First, notice that $\nu^*(A \mid T = Tx)$ is defined on a $\mathcal{T}$ -measurable subset $\mathcal{X}' \subset \mathcal{X}$ with $\nu^*(\mathcal{X}') = 1$. But as $P_\theta \ll \nu^*$ we have that $P_\theta(\mathcal{X}') = 1$, which is a requirement for conditional probability. Now observe that for $\mathcal{T}$ -measurable $C \subset \mathcal{X}'$ we have that

$$\begin{aligned} \int_C \nu^*(A \mid T = Tx) dP_\theta(x) &= \int_C \nu^*(A \mid T = Tx) h_\theta(Tx) d\nu^*(x) \\ &= \int_{\mathcal{X}} \mathbb{E}(\mathbb{1}_A \mid T = Tx) \mathbb{1}_C(x) h_\theta(Tx) d\nu^*(x) \\ &= \int_{\mathcal{X}} \mathbb{1}_A(x) \mathbb{1}_C(x) h_\theta(Tx) d\nu^*(x) \\ &= \int_{\mathcal{X}} \mathbb{1}_A(x) \mathbb{1}_C(x) dP_\theta(x) \\ &= P_\theta(A \cap C) \end{aligned}$$

which completes the proof by Proposition C.9.

□

By observing the proofs so far, we see that to check sufficiency in the Bayesian sense it is enough to check only those prior distributions that are discrete (i.e., those μ_Θ on Θ where $\mu(C) = 1$ for some at most countable subset $C \subset \Theta$).

Corollary 3.11. Assuming the nice condition of Definition 1.5 . We have that $\mathcal{T} : \mathcal{X} \rightarrow \mathcal{T}$ is sufficient (in either classical or Bayesian sense) if and only if the Bayesian sufficiency condition is satisfied for all discrete prior distributions μ_Θ on Θ .

Proof. Observe that the measures μ used in the proof of Lemma 3.9 are all discrete. Once ν^* is constructed, sufficiency follows as seen in the proof of the previous Theorem. $\square$

3.3. Fisher-Neyman factorization. There is a convenient characterization of sufficiency.

Theorem 3.12. Assume the nice condition of Definition 1.5. Then a statistic $\mathcal{T} : \mathcal{X} \rightarrow \mathcal{T}$ is sufficient if and only if the following holds: There is Borel $g : \mathcal{X} \rightarrow \mathbb{R}_{\geq 0}$ such that for all $\theta \in \Theta$ we have a Borel $h_\theta : \mathcal{T} \rightarrow \mathbb{R}_{\geq 0}$ such that

$$f(x|\theta) = h_\theta(\mathcal{T}x)g(x) \quad \text{for } P_\theta \text{ almost all } x \in X$$

Proof. First suppose this condition holds. Thus by Corollary 3.11 we only need to check Bayesian sufficiency for any discrete prior μ on Θ . Say μ is supported on $C = \{\theta_1, \theta_2, \dots\} \subset \Theta$. For $C' \subset C$ we have by Bayes' formula that

$$P(\Theta \in C' \mid X = x) = \frac{\int_{C'} f(x|\theta) d\mu(\theta)}{\int_{\Theta} f(x|\theta) d\mu(\theta)} = \frac{\sum_{\theta \in C'} h_\theta(\mathcal{T}x)g(x)\mu(\{\theta\})}{\sum_{\theta \in C} h_\theta(\mathcal{T}x)g(x)\mu(\{\theta\})} = \frac{\sum_{\theta \in C'} h_\theta(\mathcal{T}x)\mu(\{\theta\})}{\sum_{\theta \in C} h_\theta(\mathcal{T}x)\mu(\{\theta\})}.$$

Observe now that this is $\mathcal{T}$ -measurable since the sums are countable, thus Bayesian sufficiency is established.

(**Remark:** this is the reason we restrict without loss of generality to μ discrete, as otherwise it seems one cannot show measurability in the integrand unless we also assume that $(\theta, t) \mapsto h_\theta(t)$ is Borel, which we do not want to as we have not been able to verify this in Lemma 3.9).

Conversely, suppose that $\mathcal{T}$ is sufficient. Thus by Lemma 3.9 we have

$$f(x|\theta) = \frac{dP_\theta}{d\nu}(x) = \frac{dP_\theta}{d\nu^*}(x) \frac{d\nu^*}{d\nu}(x) = h_\theta(\mathcal{T}x)g(x)$$

for ν^* -almost all $x \in X$. As $P_\theta \ll \nu^*$ we have it also holds for P_θ -almost all $x \in X$. $\square$

Example 3.13. In Example 3.2 we have that

$$f(x|\theta) = P_\theta(X = x) = \theta^{T(x)}(1 - \theta)^{n - T(x)}$$

so can just take

$$h_\theta(t) = \theta^t(1 - \theta)^{n - t}$$

and $g(x) = 1$. Here $\nu(\{x\}) = 1$ is the counting measure. If instead we had a different measure $\nu(\{x\}) = q(x) > 0$ then we would have

$$f(x|\theta) = P_\theta(X = x)q(x)^{-1} = h_\theta(\mathcal{T}x)q(x)^{-1}$$

so then $g(x) = q(x)^{-1}$. This shows that $g(x)$ is just a change of density.

Example 3.14. Let $\mathcal{X} = \mathbb{R}^n$, $\Theta = \{(\mu, \sigma) \in \mathbb{R}^2 \mid \mu \in \mathbb{R}, \sigma > 0\}$ and for $\theta = (\mu, \sigma)$ let P_θ be the distribution of $(X_1, \dots, X_n)$ where each X_i is Gaussian with mean μ and variance σ^2 and the X_i are independent. Thus if ν is the Lebesgue measure on $\mathbb{R}^d$ then

$$\frac{dP_\theta}{d\nu}(x_1, \dots, x_n) = \frac{1}{\sqrt{2\pi\sigma^2}^n} \exp\left(-\frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2\right).$$

Now let $\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i$ for $x = (x_1, \dots, x_n)$ and

$$s^2 = \frac{1}{(n-1)} \sum_{i=1}^n (x_i - \bar{x})^2$$

. Thus the sum inside the exponential can be written as

$$\begin{aligned}
\sum_{i=1}^n (x_i - \bar{x} + \bar{x} - \mu) &= \sum_{i=1}^n (x_i - \bar{x})^2 + 2(\bar{x} - \mu) \sum_{i=1}^n (x_i - \bar{x}) + \sum_{i=1}^n (\bar{x} - \mu)^2. \\
&= (n-1)s^2 + 2(\bar{x} - \mu)(n\bar{x} - n\bar{x}) + \sum_{i=1}^n (\bar{x} - \mu)^2. \\
&= (n-1)s^2 + \sum_{i=1}^n (\bar{x} - \mu)^2.
\end{aligned}$$

Thus we have that

$$\frac{dP_\theta}{d\nu}(x_1, \dots, x_n) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left(-\frac{1}{2\sigma^2} \left((n-1)s^2 + \sum_{i=1}^n (\bar{x} - \mu)^2 \right) \right),$$

which may be written as $h_\theta(T(x))$ for some function h_θ where

$$T(x) = (\bar{x}, s^2).$$

Thus T , which gives the sample average and sample variance, is a sufficient statistic.

3.4. Applications to estimation. Recall that $d : \mathcal{X} \rightarrow \mathcal{A}$ is the decision function and $L : \Theta \times \mathcal{A} \rightarrow [0, \infty)$ denotes the loss function. The risk is then

$$R(\theta, d) = \int_{\mathcal{X}} L(\theta, d(x)) dP_\theta(x).$$

Recall (see Theorem 3.10) that if $T : \mathcal{X} \rightarrow \mathcal{T}$ is a sufficient statistic then there are families of measures ν_t^* , for $t \in \mathcal{T}$, such that

$$P_\theta(B \mid T = t) = \nu_t^*(B) \quad \text{for all } B \subset \mathcal{X}.$$

Theorem 3.15 (Rao-Blackwell theorem). Suppose that $\mathcal{A} = \mathbb{R}^d$ and that for each fixed θ , the map $a \mapsto L(\theta, a)$ is convex. Suppose also that $x \mapsto d(x)$ is an L^1 function for each P_θ , in other words

$$\int_{\mathcal{X}} \|d(x)\| dP_\theta < \infty.$$

Suppose that $T : \mathcal{X} \rightarrow \mathcal{T}$ is a sufficient statistic for the family of P_θ . Then there is a well-defined decision rule given by

$$\hat{d}(x) = E_\theta(d \mid T = Tx) = \int_{\mathcal{X}} d(x') d\nu_{Tx}^*(x')$$

which satisfies that

$$R(\theta, \hat{d}) \leq R(\theta, d).$$

Proof. Note that as d is L^1 with respect to P_θ , we must have that it is L^1 with respect to $P_\theta(\cdot \mid T = Tx) = \nu_t^*(\cdot)$ for P_θ -almost all x . Thus the expression for $\hat{d}(x)$ is well defined. By Jensen's theorem, we have that

$$L(\theta, \hat{d}(x)) = L \left(\theta, \int_{\mathcal{X}} d(x') d\nu_{Tx}^*(x') \right) \leq \int_{\mathcal{X}} L(\theta, d(x')) d\nu_{Tx}^*(x')$$

Now we integrate both sides with respect to P_θ to get that

$$\begin{aligned}
R(\theta, \hat{d}) &\leq \int_{\mathcal{X}} \int_{\mathcal{X}} L(\theta, d(x')) d\nu_{Tx}^*(x') dP_\theta(x) \\
&= \int_{\mathcal{X}} L(\theta, d(x)) dP_\theta(x) \\
&= R(\theta, d)
\end{aligned}$$

where we used the law of total probability applied to $P_\theta(\cdot \mid T = Tx) = \nu_{Tx}^*$. $\square$

Thus the Rao-Blackwell theorem allows us to construct more efficient estimators that are a function of the sufficient statistic T .

Examples of such convex loss functions include the square loss

$$L(\theta, a) = |\theta - a|^2.$$

APPENDIX A. REVIEW OF PROBABILITY

APPENDIX B. BOREL SPACES

Lemma B.1. Let X be a second-countable metric space. Then there exists an embedding $\iota : X \hookrightarrow Y$ where Y is compact and ι is a homeomorphism onto its image $\iota(X)$.

Proof. Let $u_1, u_2, \dots$ be a countable dense subset of X . Suppose that $d(x, x') \leq 1$ for all x, x' . Define

$$\iota : X \rightarrow [0, 1]^\mathbb{N}$$

to be

$$\iota(x) = (d(x, u_i))_{i=1}^\infty$$

where $Y = [0, 1]^\mathbb{N}$ has the product topology, thus is a compact metric space. Continuity and injectivity of ι is clear. Now suppose that $x_1, x_2, \dots \in X$ and $x \in X$ are such that $\iota(x_i) \rightarrow \iota(x)$ as $i \rightarrow \infty$. Fix $\epsilon > 0$ and choose u_j so that $d(x, u_j) < \epsilon$. Now for all sufficiently large i we have that $|d(x_i, u_j) - d(x, u_j)| < \epsilon$. This implies that $d(x_i, u_j) < 2\epsilon$. Now

$$d(x, x_i) \leq d(x, u_j) + d(u_j, x_i) \leq \epsilon + 2\epsilon = 3\epsilon.$$

Thus ι^{-1} is continuous on $\iota(X)$. $\square$

Lemma B.2. Let X be a complete metric space, with metric d_X , and suppose that $Y \subset X$ is a subset of X such that the subspace topology on Y is induced by a complete metric d_Y on Y . Then Y is a countable intersection of open subsets of X , and is thus a Borel subset of Y .

Proof. We first show that for each $\epsilon > 0$ and $y \in Y$ there exists an open set $V_{y,\epsilon}$ of X such that $y \in V$ and $\text{diam}_{d_X}(V_{y,\epsilon}) \leq \epsilon$ and $\text{diam}_{d_Y}(V_{y,\epsilon} \cap Y) \leq \epsilon$. To see this, observe since d_X and d_Y induce the same topology on Y , there is an open set $\tilde{V}$ of X such that $\tilde{V} \cap Y$ is the ball of radius $\epsilon/2$ around y in the d_Y metric. Hence we can take

$$V_{y,\epsilon} = \tilde{V} \cap B_{d_X}(y, \epsilon/2).$$

Now let $U_n = \bigcup_{y \in Y} V_{y, 1/n}$ (Remark: to avoid axiom of choice, just take the union of y and all the different sets of the form $\tilde{V}$). Observe now that $Y \subset U_n$ for all n . We claim that

$$Y = \bigcap_{n=1}^\infty U_n.$$

Thus suppose that $x \in U_n$ for all $n \geq 1$. We must show $x \in Y$. We will show that, by induction on n , that there exists a sequence of open subsets $W_1 \supset W_2 \supset \dots$ of X and $y_n \in Y \cap W_n$ such that $\text{diam}_{d_X}(W_n) \leq \frac{1}{n}$, $\text{diam}_{d_Y}(W_n \cap Y) \leq \frac{1}{n}$, $x \in W_n$. We can take $W_1 = V_{y_1, 1}$ for some $y_1 \in Y$ since $x \in U_1$. Now suppose W_n has been constructed. Thus there exists $\epsilon > 0$ so that $B_{d_X}(x, \epsilon) \subset W_n$ as W_n is open and $x \in W_n$. Now take $N > n$ such that $1/N < \epsilon$. Now since $x \in U_N$ there exists $y_{n+1} \in Y$ and $W_{n+1} \subset X$

open so that $y_{n+1}, x \in W_{n+1}$ and $\text{diam}_{d_X}(W_{n+1}) \leq 1/N$, $\text{diam}_{d_Y}(W_{n+1} \cap Y) \leq 1/N$. Now observe that $W_{n+1} \subset B_{d_X}(x, \epsilon) \subset W_n$ since if $x' \in W_{n+1}$ then

$$d_X(x, x') \leq \text{diam}_{d_X}(W_{n+1}) \leq 1/N < \epsilon.$$

Thus in summary $y_{n+1} \in W_{n+1}$ is non-empty as desired and $W_{n+1} \subset W_n$ and

$$\text{diam}_{d_Y}(W_{n+1} \cap Y) \leq \frac{1}{N} \leq \frac{1}{n+1}.$$

It now follows that $y_n \rightarrow x$ in d_X metric and that $y_1, y_2, \dots$ is a Cauchy sequence in Y . Thus by completeness, we have that $y_n \rightarrow y \in Y$ for some $y \in Y$. However

$$d_X(x, y) \leq d_X(x, y_n) + d_X(y_n, y) \rightarrow 0$$

as $n \rightarrow \infty$ since $d_Y(y, y_n) \rightarrow 0$ and d_X and d_Y induce the same topology. Thus $x = y \in Y$ as required. $\square$

Corollary B.3. Let X be a second countable complete metric space. Then X is homeomorphic to a Borel subset of the compact metric space $[0, 1]^{\mathbb{N}}$. In fact, we can take this Borel subset to be a countable intersection of open subsets.

Definition B.4. A *standard Borel space* is a measurable space $(X, \mathcal{B})$ where X is a second-countable complete metric space and $\mathcal{B}$ is the Borel σ -algebra.

APPENDIX C. CONDITIONAL EXPECTATION AND PROBABILITY

C.1. Conditional expectation.

Theorem C.1. Let $(X, \mathcal{B}, \mu)$ be a probability space and let $\mathcal{A} \subset \mathcal{B}$ be a sub- σ -algebra of $\mathcal{B}$. Then for $f \in L^1(X, \mathcal{B}, \mu)$ there exists a unique $g \in L^1(X, \mathcal{A}, \mu)$ such that

$$\int_A f d\mu = \int_A g d\mu \quad \text{for all } A \in \mathcal{A}.$$

Definition C.2. We say that the $g \in L^1(X, \mathcal{B}, \mu)$ the *conditional expectation of f w.r.t. $\mathcal{A}$* , and denote it by $g = E(f|\mathcal{A})$.

Proof. If $f \in L^2(X, \mathcal{B}, \mu)$ then we define $E(f|\mathcal{A})$ to be the orthogonal projection of f onto the closed subspace $L^2(X, \mathcal{A}, \mu)$. Indeed, for $A \in \mathcal{A}$ we have that $\mathbf{1}_A \in L^2(X, \mathcal{A}, \mu)$ and so

$$\int_A f d\mu = \langle f, \mathbf{1}_A \rangle = \langle E(f|\mathcal{A}), \mathbf{1}_A \rangle = \int_A E(f|\mathcal{A}) d\mu$$

as required. We will now extend $E(\cdot|\mathcal{A})$ to L^1 . By density of L^2 in L^1 it is enough to show that $E(\cdot|\mathcal{A})$ is a bounded operator w.r.t. to the L^1 norm on L^2 . For $f \in L^2(X, \mathcal{B}, \mu)$ let $\psi \in L^\infty(X, \mathcal{A}, \mu)$ be such that $|E(f|\mathcal{A})| = \psi E(f|\mathcal{A})$ and $\|\psi\|_\infty \leq 1$. Thus

$$\|E(f|\mathcal{A})\|_1 = \int_X \psi E(f|\mathcal{A}) d\mu = \langle E(f|\mathcal{A}), \bar{\psi} \rangle = \langle f, \bar{\psi} \rangle = \int_X \psi f d\mu \leq \|f\|_1.$$

Thus $E(\cdot|\mathcal{A})$ has L^1 operator norm at most 1 and thus can be extended to $L^1(X, \mathcal{B}, \mu) \rightarrow L^1(X, \mathcal{A}, \mu)$.

For uniqueness, suppose that $g_1, g_2 \in L^1(X, \mathcal{A}, \mu)$ are real valued functions such that

$$\int_A g_1 d\mu = \int_A g_2 d\mu \quad \text{for all } A \in \mathcal{A}.$$

Let

$$A_\epsilon = \{x \in X \mid g_1(x) \leq g_2(x) - \epsilon\}$$

and observe that $A_\epsilon \in \mathcal{A}$. Thus

$$\int_A g_2 = \int_A g_1 d\mu \leq \int_A g_2 d\mu - \epsilon \mu(A)$$

which implies that $\mu(A_\epsilon) = 0$. Thus $g_1 \geq g_2$ almost everywhere. By symmetry, $g_1 = g_2$ almost everywhere. For g_1, g_2 complex valued uniqueness also follows by equating real and imaginary parts. $\square$

Definition C.3. Let $T : X \rightarrow Y$ be a measurable map between measure spaces $(X, \mathcal{B}_X)$ and $(Y, \mathcal{B}_Y)$. We define $T^{-1}\mathcal{B}_Y = \{T^{-1}B \mid B \in \mathcal{B}_Y\}$.

Proposition C.4. Suppose that $(X, \mathcal{B}_X, \mu_X)$ is a probability space, $(Y, \mathcal{B}_Y)$ is a measurable space $T : X \rightarrow Y$ is measurable. Then any function $f : X \rightarrow \mathbb{R}$ that is $T^{-1}\mathcal{B}_Y$ measurable (that is, $f^{-1}B \in T^{-1}\mathcal{B}_Y$ for all Borel $B \subset \mathbb{R}$) can be written as $f = g \circ T$ for some measurable $g : Y \rightarrow \mathbb{R}$.

In other words, if we let $\mu_Y = T^*\mu_X$ be the push-forward measure then we have an isometric embedding

$$L^p(Y, \mathcal{B}_Y, \mu_Y) \hookrightarrow L^p(X, \mathcal{B}_X, \mu_X)$$

given by $g \mapsto g \circ T$ whose image is exactly $L^p(\mathcal{X}, T^{-1}\mathcal{B}_Y, \mu_X)$.

Proof. If $f = \mathbb{1}_A$ for some $A \in T^{-1}\mathcal{B}_Y$ then $A = T^{-1}B$ for some $B \in \mathcal{B}_Y$ and so we can take $g = \mathbb{1}_B$. For $f \geq 0$ note that we can write

$$f = \sum_{i=1}^{\infty} c_i \mathbb{1}_{A_i}$$

for some $c_i \geq 0$ and $A_i \in T^{-1}\mathcal{B}_Y$ where the limit is pointwise and so $\mathbb{1}_{A_i} = \mathbb{1}_{B_i} \circ T$ for some $B_i \in \mathcal{B}_Y$. Thus $f = g \circ T$ where

$$g(y) = \sum_{i=1}^{\infty} c_i \mathbb{1}_{B_i}(y)$$

converges pointwise as it is monotonic $c_i > 0$. For arbitrary measurable f , note that we can write it as a linear combination of non-negative functions.

Now note that for $g \geq 0$ we have that

$$\int_X g \circ T d\mu_X = \int_Y g d\mu_Y$$

which can be shown as follows. Writing

$$g(y) = \sum_{i=1}^{\infty} c_i \mathbb{1}_{B_i}(y)$$

and using the monotone convergence theorem we have that

$$\int_X g(Tx) d\mu_X(x) = \sum_{i=1}^{\infty} c_i \int_X \mathbb{1}_{B_i}(Tx) d\mu_X(x) = \sum_{i=1}^{\infty} c_i \mu_X(T^{-1}B_i) = \sum_{i=1}^{\infty} c_i \mu_Y(B_i) = \int_Y g(y) d\mu_Y(y).$$

This shows that for $1 \leq p < \infty$ we have that

$$\|g \circ T\|_p^p = \int_X |g|^p(Tx) d\mu_X(x) = \int_Y |g|^p(y) d\mu_Y(y) = \|g\|_p^p$$

and so indeed the embedding is an isometry. $\square$

This proposition motivates us to define conditional expectation with respect to a random variable (or factor map).

Definition C.5. Let $(X, \mathcal{B}, \mu)$ be a probability space and suppose that $T : X \rightarrow Y$ is measurable where $(Y, \mathcal{B}_Y)$ is a measure space (i.e., a random variable). Then for $f \in L^p(\mathcal{X}, \mathcal{B}, \mu)$ (for $p = 1, 2, \infty$) we define a function on $Y \rightarrow \mathbb{C}$ denoted $y \mapsto E(f|T = y)$ that is characterized for μ -almost all x by

$$E(f|\mathcal{T}^{-1}\mathcal{B}_Y)(x) = E(f|T = y) \quad \text{where } y = Tx.$$

In other words, the function $E(f|\mathcal{T}^{-1}\mathcal{B}_Y) \in L^p(\mathcal{X}, \mathcal{T}^{-1}\mathcal{B}_Y, \mu)$ is identified with a function in $L^p(Y, \mathcal{B}_Y, T^*\mu)$ using the isometry in the proposition above and this function at $y \in Y$ is denoted $E(f|T = y)$. Note that $E(f|T = y)$ is uniquely defined for $T^*\mu_X$ -almost all $y \in Y$.

Note that $E(f|T = y)$ can be characterized by the property that

$$\int_A E(f|T = y) d\mu_Y(y) = \int_{T^{-1}A} f(x) d\mu_X(x) \quad \text{for all } A \in \mathcal{B}_Y.$$

C.2. Conditional Probability.

Theorem C.6 (Conditional probability). Let X be a second-countable complete metric space with Borel σ -algebra $\mathcal{B}$ and suppose $\mathcal{A} \subset \mathcal{B}$ is a sub σ -algebra. Let μ be a probability measure on $(X, \mathcal{B})$. Then there exists $X' \in \mathcal{A}$ with $\mu(X') = 1$ and a family of probability measures $\mu_x^{\mathcal{A}}$ for $x \in X'$ such that the following holds.

- For $\mathcal{B}$ -measurable functions $f : X \rightarrow \mathbb{C}$ such that $\|f\|_\infty < \infty$, we have that the mapping $X' \rightarrow \mathbb{C}$ given by

$$x \mapsto \int_X f(u) d\mu_x^{\mathcal{A}}(u)$$

is $\mathcal{A}$ -measurable. The same is true if we replace $\|f\|_\infty < \infty$ with $f(x) \in \mathbb{R}_{\geq 0}$ for all $x \in X$.

- For $f : X \rightarrow \mathbb{C}$ with $\|f\|_1 < \infty$ we have that¹

$$E(f|\mathcal{A})(x) = \int_X f(u) d\mu_x^{\mathcal{A}}(u) \quad \text{for almost all } x \in X.$$

Definition C.7. We call such a family of measures $\mu_x^{\mathcal{A}}$ *conditional probabilities of μ w.r.t $\mathcal{A}$* .

Proof. Using Corollary B.3 we may assume X is compact. Thus there exists a countable dense $\mathbb{Q}$ -vector subspace V of $C(X)$ such that $1 \in V$. Now for each $f \in V$ choose a genuine function (not an equivalence class of functions) $g_f : X \rightarrow \mathbb{C}$ such that

$$g_f(x) = E(f|\mathcal{A})(x)$$

and

$$\|g_f\|_\infty \leq \|f\|_\infty.$$

Now for any $\alpha, \beta \in \mathbb{Q}$ and $f_1, f_2 \in V$ we have that

$$g_{\alpha f_1 + \beta f_2}(x) = \alpha g_{f_1}(x) + \beta g_{f_2}(x) \quad \text{for almost all } x \in X.$$

We also have $g_1(x) = 1$ and $g_f(x) \geq 0$ for positive $f \in V$ (i.e., $f \geq 0$) for almost all $x \in X$. Thus we may choose $X' \subset X$ to be a set with $\mu(X') = 1$ on which all of these properties hold (as there are only countably many equations). Thus for each $x \in X'$ we have a linear function $\Lambda'_x : V \rightarrow \mathbb{C}$ given by

$$\Lambda'_x(f) = g_f(x).$$

¹The right hand side may be undefined on a null set of x .

Note that Λ'_x is uniformly continuous (since $\|g_f\|_\infty \leq \|f\|_\infty$) and thus it extends to a linear functional $\Lambda_x : C(X) = \overline{V} \rightarrow \mathbb{C}$ which also satisfies that

$$|\Lambda_x(f)| \leq \|f\|_\infty.$$

We wish to apply the Riesz representation theorem, thus it remains to show that Λ_x is a positive linear functional. Now suppose $f \in C(X)$ is positive. Then there exists $f_i \in V$ such that $f_i \rightarrow f$ uniformly. Now fix a positive integer q and let $\epsilon = 1/q$. Then $\Lambda_x(\epsilon) = \epsilon$ since $g_1 = 1$. For i sufficiently large we have that we have $\|f_i - f\|_\infty < \epsilon$ and so $f_i + \epsilon \geq 0$ as f is positive. As $f_i + \epsilon \in V$ we have that

$$\Lambda_x(f_i) + \epsilon = g_{f_i + \epsilon}(x) \geq 0.$$

Thus

$$\Lambda_x(f) \geq \Lambda_x(f_i) - \Lambda_x(|f_i - f|) \geq \Lambda_x(f_i) - \epsilon > -2\epsilon.$$

This shows that $\Lambda_x(f) \geq 0$ and thus Λ_x is a positive linear continuous functional on $C(X)$ with $\Lambda_x(1) = 1$. Thus there exists a measure μ_x such that

$$\Lambda_x(f) = \int_X f(u) d\mu_x(u) \quad \text{for all } f \in C(X), x \in X'.$$

In particular this means that

$$E(f|\mathcal{A})(x) = \int_X f(u) d\mu_x(u) \quad \text{for all } f \in V.$$

This means that

$$(4) \quad x \mapsto \int_{X'} f(u) d\mu_x(u) \text{ is } \mathcal{A}\text{-measurable on } X'$$

and

$$(5) \quad \int_A \int_{X'} f(u) d\mu_x(u) d\mu(x) = \int_A f(x) d\mu(x) \quad \text{for all } A \in \mathcal{A}$$

holds for all $f \in V$. By density of V in $C(X)$, this also holds for all $f \in C(X)$ and the integrand is indeed $\mathcal{A}$ measurable (it is a pointwise limit of $\mathcal{A}$ measurable functions). Since the indicator function of an open set is a monotonic limit of continuous functions, we have that (4) and (5) holds for $f = \mathbb{1}_U$ for all $U \subset X$ open. It also holds for closed sets by linearity. Finally, observe that it also holds for an intersection $U \cap C$ where U is open and C is closed since such a set is a countable nested intersection of open sets.

Now let

$$\mathcal{R} = \{\cup_{i=1}^n U_i \cap C_i \mid U_1, \dots, U_n \text{ are open and } C_1, \dots, C_n \text{ are closed.}\}$$

be the set of all finite unions of sets that can be written as an intersection of an open set and a closed set. We claim that any element in $\mathcal{R}$ can be written as a finite disjoint union of an intersection of an open set with a closed set. Indeed, any such finite union belongs to a finite σ -algebra generated by such sets but any atom of this σ -algebra must take this form. It now follows that (4) and (5) holds for $f = \mathbb{1}_R$ for all $R \in \mathcal{R}$. By the monotone convergence theorem, these also hold on the monotone class generated by $\mathcal{R}$. Thus by the monotone class lemma, it holds on all Borel sets since $\mathcal{R}$ is an algebra. Any measurable $f : X \rightarrow \mathbb{R}_{\geq 0}$ is a monotonic limit of simple functions and so it holds for such f too, which completes the proof for such f . The theorem now follows for general $f : X \rightarrow \mathbb{C}$ with $\|f\|_1$ as desired by linearity. $\square$

We now show that the conditional probabilities are almost surely unique and provide some characterizations that are practical to use. In the following proposition, we fix $X, \mathcal{A}, \mu$ as in Theorem C.6.

Proposition C.8. If μ_x^A and ν_x^A are two conditional probabilities, for some $x \in X' \in \mathcal{A}$ of full measure, then $\mu_x^A = \nu_x^A$ for μ -almost all $x \in X$.

Proof. We for each $f \in C(X)$ we have that

$$\int_X f(u) d\mu_x^A(u) = E(f|\mathcal{A})(x) = \int_X f(u) d\nu_x^A(u)$$

holds for all $x \in X_f \in \mathcal{A}$ with $\mu(X_f) = 1$. Applying this to a countable dense set of $f \in C(X)$ we can find a single set $X' \in \mathcal{A}$ with $\mu(X') = 1$ so that this identity holds for all continuous functions $f \in C(X)$. It follows by uniqueness in Riesz-representation theorem that $\mu_x^A = \nu_x^A$ for all $x \in X'$. $\square$

Proposition C.9. Suppose there is a set $X' \in \mathcal{A}$ of full measure and a family of probability measures μ_x for $x \in X'$. Then μ_x form conditional probabilities of μ with respect to $\mathcal{A}$ if and only if the following holds:

For all Borel $B \subset X$ we have

$$x \mapsto \int_X \mu_x(B) d\mu(x)$$

is $\mathcal{A}$ measurable on X' and for all $A \in \mathcal{A}$ we have that

$$\int_A \mu_x(B) d\mu(x) = \mu(A \cap B).$$

Proof. This equation may be written as

$$\int_A \int_X f(u) d\mu_x(u) d\mu(x) = \int_A f(x) d\mu(x)$$

where $f = 1_B$. Thus the μ_x are conditional probabilities if and only if this holds for all continuous f and thus if and only if it holds for all positive functions f . But by the monotone convergence theorem, we have that if it holds for all indicator functions it holds for all positive functions. $\square$

C.3. Conditional probability with respect to a random variable. In this subsection, $\mathcal{A}$ is a countable generated σ -algebra of $\mathcal{B}$ where $(X, \mathcal{B})$ is standard Borel. So $\mathcal{A} = \sigma(\{A_1, A_2, \dots\})$ where $A_1, A_2, \dots \in \mathcal{A}$. For each $x \in X$ we can define the atom

$$[x]_{\mathcal{A}} = \bigcap_{i, x \in A_i} A_i \cap \bigcap_{j, x \notin A_j} (X \setminus A_j)$$

containing x .

Proposition C.10. For $x \in X$ we have that $[x]$ is the smallest subset of $\mathcal{A}$ containing x .

Proof. We must show that for each $A \in \mathcal{A}$ we have either $[x] \subset A$ or $[x] \cap A = \emptyset$. This is true for each A_i and so one must show that the sets in $\mathcal{A}$ satisfying this property forms a σ -algebra. $\square$

Proposition C.11. There is a set $X'' \subset X$ with $X'' \in \mathcal{A}$ and $\mu(X'') = 1$ such that for any $x_1, x_2 \in X''$ we have that $\mu_{x_1}^A = \mu_{x_2}^A$ if and only if $[x_1] = [x_2]$. Moreover, $\mu_x^A([x]) = 1$ for $x \in X''$.

Proof. Let $\mu_x = \mu_x^A$ for $x \in X'$ be a set of conditional probabilities, where $X' \in \mathcal{A}$ and $\mu(X') = 1$. For almost all $x \in X'$ we have that

$$\mu_x(A_i) = \int \mathbb{1}_{A_i} d\mu_x = E(\mathbb{1}_{A_i}|\mathcal{A})(x) = \mathbb{1}_{A_i}(x).$$

Say this occurs on $X'' \subset X'$ with $\mu(X'') = 1$. Note that since all terms are $\mathcal{A}$ -measurable, we have $X'' \in \mathcal{A}$. Now let $x_1, x_2 \in X''$. If $\mu_{x_1} = \mu_{x_2}$ then $\mathbb{1}_{A_i}(x_1) = \mathbb{1}_{A_i}(x_2)$ for all i and thus $[x_1] = [x_2]$. Conversely,

suppose that $[x_1] = [x_2]$. Thus any $\mathcal{A}$ -measurable function must have equal values on x_1 and x_2 . But for any $f \in C(X)$ we have that

$$x \mapsto \int f d\mu_x$$

is $\mathcal{A}$ -measurable. Thus

$$\int f d\mu_{x_1} = \int f d\mu_{x_2}$$

holds for all $f \in C(X)$ and thus $\mu_{x_1} = \mu_{x_2}$. $\square$

Theorem C.12 (Law of total probability). Suppose that $T : X \rightarrow Y$ is a measurable map between standard Borel spaces and let $\mathcal{B}_T := T^{-1}\mathcal{B}_Y$ be the induced σ -algebra. There exists Borel $Y' \subset Y$ and probability measures μ_y on X , for each $y \in Y'$, such that the following hold.

- $\mu(T^{-1}Y') = 1$
- For all $y \in Y'$, the measure μ_y is supported on $T^{-1}(\{y\})$.
- For all $x \in T^{-1}(Y')$ we have that $\mu_{Tx} = \mu_x^{\mathcal{B}_T}$. That is the family μ_{Tx} for $x \in T^{-1}(Y')$ is a family of conditional probability measures for the σ -algebra $\mathcal{B}_T$.

Moreover, for bounded Borel $f : X \rightarrow \mathbb{C}$ and $B \subset Y$ Borel we have that

$$\int_{T^{-1}B} f(x) d\mu(x) = \int_B \left(\int_X f(u) d\mu_y(u) \right) d(T^*\mu)(y).$$

Proof. The induced σ -algebra $\mathcal{B}_T := T^{-1}\mathcal{B}_Y$ is countably generated and its atoms are precisely the fibers

$$T^{-1}(\{y\}) \quad \text{for } y \in Y.$$

In other words,

$$[x]_{\mathcal{B}_T} = T^{-1}(\{Tx\}) \quad \text{for } x \in X.$$

We thus have by the above proposition that the conditional measure $\mu_x^{\mathcal{B}_T}$ is supported on the fiber $T^{-1}(\{Tx\})$ and only depends on the fiber. More precisely, we have a family of measures μ_y on X , $y \in Y'$ for some Borel $Y' \subset Y$ such that μ_y is supported on $T^{-1}(\{y\})$ and such that for bounded Borel $f : X \rightarrow \mathbb{C}$ we have that

$$x \mapsto \int_X f(u) d\mu_{Tx}(u)$$

is $\mathcal{B}_T$ measurable and

$$\int_A f(x) d\mu(x) = \int_A \int f(u) d\mu_{Tx}(u) d\mu(x) \quad \text{for } A \in \mathcal{B}_T.$$

Now by Proposition C.4 we have that there exists $g : Y' \rightarrow \mathbb{C}$ which is Borel so that

$$g(Tx) = \int_X f(u) d\mu_{Tx}(u).$$

In particular, this means that

$$y \mapsto g(y) = \int_X f(u) d\mu_y(u)$$

is Borel on Y' .

Now for Borel $B \subset Y'$ we have that $A = T^{-1}B \in \mathcal{B}_T$ and the equation above says that

$$\begin{aligned} \int_A f(x) d\mu(x) &= \int_A g(Tx) d\mu(x) \\ &= \int_B g(y) d(T^*\mu)(y) \\ &= \int_B \left(\int_X f(u) d\mu_y(u) \right) d(T^*\mu)(y) \end{aligned}$$

which completes the proof. $\square$

It is customary to define

$$\mu(B \mid T = y) := \mu_y(B)$$

and

$$E(f \mid T = y) := E(f \mid \mathcal{B}_T)(x)$$

for $y = Tx$. The law of total probability then says

$$\mu(A \cap T^{-1}B) = \int_B \mu(A \mid T = y) d\mu_Y(y)$$

for any Borel $A \subset X, B \subset Y$ where $\mu_Y = T^*\mu$.

APPENDIX D. CONDITIONAL INDEPENDENCE

In this section, fix a standard Borel probability space (Ω, μ) and suppose we have maps $X : \Omega \rightarrow \mathcal{X}$, $Y : \Omega \rightarrow \mathcal{Y}$ and $T : \Omega \rightarrow \mathcal{T}$ to some standard Borel spaces. We also let $\mathcal{A}_X, \mathcal{A}_Y, \mathcal{A}_T$ be the σ -algebras induced by these maps, respectively.

Definition D.1. We say that X and Y are conditionally independent given T if for all $A \subset \mathcal{X}, B \subset \mathcal{Y}$ we have that

$$\mu(X^{-1}A \cap Y^{-1}B \mid T = t) = \mu(X^{-1}A \mid T = t) \mu(Y^{-1}B \mid T = t) \quad \text{for } T^*\mu\text{-almost all } t \in \mathcal{T}.$$

In particular, if T is constant then we say that X and Y are independent, i.e., X and Y are independent if and only if for all $A \subset \mathcal{X}, B \subset \mathcal{Y}$ we have

$$\mu(X^{-1}A \cap Y^{-1}B) = \mu(X^{-1}A) \mu(Y^{-1}B).$$

Example D.2. Suppose $\Omega = [0, 1] \times [0, 1] \cup [1, 2] \times [1, 2]$ is the disjoint union of two squares and μ is the uniform probability measure on Ω . Let $X, Y : \Omega \rightarrow [0, 2]$ denote the projections $X(x, y) = x, Y(x, y) = y$ and let $T : \Omega \rightarrow \{0, 1\}$ be the map $T(x, y) = \mathbf{1}_{[0, 1]^2}((x, y))$. We will show that X, Y are conditionally independent given T but are not independent. First, note that for $U \subset \Omega$ we have that

$$\mu(U) = \frac{1}{2} \lambda_2(U)$$

where λ_d denotes the Lebesgue measure on $\mathbb{R}^d$. It now follows that

$$\mu(U \mid T = 0) = \lambda_2(U \cap [0, 1]^2)$$

and

$$\mu(U \mid T = 1) = \lambda_2(U \cap [1, 2]^2).$$

Now for $A, B \subset [0, 2]$ we have that

$$\begin{aligned}
\mu(X^{-1}A \cap Y^{-1}B | T = 0) &= \lambda_2(A \times B \cap [0, 1]^2) \\
&= \lambda_1(A \cap [0, 1])\lambda_1(B \cap [0, 1]) \\
&= \lambda_2(X^{-1}A \cap [0, 1]^2)\lambda_2(Y^{-1}B \cap [0, 1]^2) \\
&= \mu(X^{-1}A | T = 0)\mu(Y^{-1}B | T = 0)
\end{aligned}$$

and a similar calculation holds for conditioning on $T = 1$, which shows conditional independence.

Independence does not since if we take $A = [0, 1], B = [1, 2]$ then

$$\mu(X^{-1}A \times Y^{-1}B) = \mu([0, 1]^2 \cap [1, 2]^2) = 0$$

but

$$\mu(X^{-1}A) = \mu([0, 1]^2) = \frac{1}{2}$$

and

$$\mu(Y^{-1}B) = \mu([1, 2]^2) = \frac{1}{2}.$$

Example D.3 (Independent but not conditionally). Let $\Omega = [0, 1]^2$ and suppose μ is the uniform measure on Ω . Let $X, Y : \Omega \rightarrow [0, 1]$ be the projections. Then X and Y are independent. Let $T : \Omega \rightarrow [0, 2]$ by $T(x, y) = x + y$. Then in fact X and Y are not conditionally independent given T . For example, let $A = B = [0, 1/2]$. Then for $1 < t < \frac{3}{2}$ we have that $(A \times B) \cap T^{-1}(t) = \emptyset$ and so $P(A \times B | T = t) = 0$. However $P(X^{-1}A | T = t) > 0$ and $P(Y^{-1}B | T = t) > 0$.

Observe that we have a map $S : \Omega \rightarrow \mathcal{Y} \times \mathcal{T}$ given by $S(\omega) = (Y(\omega), T(\omega))$. We let

$$\mu(A | Y = y, T = t) = \mu(A | S = (y, t)).$$

Lemma D.4. We have that X and Y are conditionally independent given T if and only if for all $A \subset \mathcal{X}$ we have that

$$\mu(X^{-1}A | Y = Y\omega, T = T\omega) = \mu(X^{-1}A | T = T\omega) \quad \text{for } \mu\text{-almost all } \omega \in \Omega.$$

Proof. First assume X and Y are conditionally independent. Thus we must show that the right hand side satisfies the definition of conditional probability with respect to the random variable S . The right hand side is S measurable, as required. Thus it remains to show that

$$\int_{S^{-1}U} \mu(X^{-1}A | T = T\omega) d\mu(\omega) = \mu(X^{-1}A \cap S^{-1}U)$$

for all $U \subset \mathcal{Y} \times \mathcal{T}$ and $A \subset \mathcal{X}$. By the monotone class lemma, it is enough to show it for $U = B \times C$. We have

$$\begin{aligned}
\int_{S^{-1}(B \times C)} \mu(X^{-1}A \mid T = T\omega) d\mu(\omega) &= \int_{Y^{-1}B \cap T^{-1}C} \mu(X^{-1}A \mid T = T\omega) d\mu(\omega) \\
&= \int_{\Omega} \mathbb{1}_{Y^{-1}B}(\omega) \mathbb{1}_{T^{-1}C}(\omega) \mu(X^{-1}A \mid T = T\omega) d\mu(\omega) \\
&= \int_{\Omega} E(\mathbb{1}_{Y^{-1}B} \mid T = T\omega) \mathbb{1}_{T^{-1}C}(\omega) \mu(X^{-1}A \mid T = T\omega) d\mu(\omega) \\
&= \int_{T^{-1}C} \mu(Y^{-1}B \mid T = T\omega) \mu(X^{-1}A \mid T = T\omega) d\mu(\omega) \\
&= \int_{T^{-1}C} \mu(Y^{-1}B \cap X^{-1}A \mid T = T\omega) d\mu(\omega) \\
&= \mu(Y^{-1}B \cap X^{-1}A \cap T^{-1}C) \\
&= \mu(X^{-1}A \cap S^{-1}(B \times C))
\end{aligned}$$

as required.

Now we prove the converse. Thus we must assume that for $A \subset \mathcal{X}$ we have

$$\mu(X^{-1}A \mid Y = Y\omega, T = T\omega) = \mu(X^{-1}A \mid T = T\omega) \quad \text{for } \mu\text{-almost all } \omega \in \Omega$$

and show that

$$\int_C \mu(X^{-1}A \mid T = t) \mu(Y^{-1}B \mid T = t) dT^* \mu(t) = \int_C \mu(X^{-1}A \cap Y^{-1}B \mid T = t) dT^* \mu(t)$$

for all $C \subset \mathcal{T}$. Using the law of total probability we have

$$\begin{aligned}
\int_C \mu(X^{-1}A \cap Y^{-1}B \mid T = t) dT^* \mu(t) &= \mu(X^{-1}A \cap Y^{-1}B \cap T^{-1}C) \\
&= \mu(X^{-1}A \cap S^{-1}(B \times C)) \\
&= \int_{B \times C} \mu(X^{-1}A \mid Y = y, T = t) d(S^* \mu)(y, t) \\
&= \int_{Y^{-1}B \cap T^{-1}C} \mu(X^{-1}A \mid Y = Y\omega, T = T\omega) d\mu(\omega) \\
&= \int_{Y^{-1}B \cap T^{-1}C} \mu(X^{-1}A \mid T = T\omega) d\mu(\omega) \\
&= \int_{\Omega} \mathbb{1}_{Y^{-1}B}(\omega) \mathbb{1}_{T^{-1}C}(\omega) \mu(X^{-1}A \mid T = T\omega) d\mu(\omega) \\
&= \int_{\Omega} E(\mathbb{1}_{Y^{-1}B} \mid T = T\omega) \mathbb{1}_{T^{-1}C}(\omega) \mu(X^{-1}A \mid T = T\omega) d\mu(\omega) \\
&= \int_{\Omega} \mu(Y^{-1}B \mid T = T\omega) \mathbb{1}_{T^{-1}C}(\omega) \mu(X^{-1}A \mid T = T\omega) d\mu(\omega) \\
&= \int_C \mu(Y^{-1}B \mid T = t) \mu(X^{-1}A \mid T = t) dT^* \mu(t)
\end{aligned}$$

□

APPENDIX E. USEFUL RESULTS

Lemma E.1. Suppose that ν is a σ -finite positive measure on X . Then there exists a probability measure ν' such that $\nu \ll \nu'$ and $\nu' \ll \nu$ (i.e., $\nu(A) = 0$ if and only if $\nu'(A) = 0$).

Proof. We only need to consider the case $\nu(X) = \infty$. There exists a countable partition $X_1, \dots$ of X such that $0 < \nu(X_i) = d_i < \infty$. Now define

$$\nu'(A) = \sum_{i=1}^{\infty} 2^{-i} d_i^{-1} \nu(A \cap X_i).$$

□

Lemma E.2. Let ν be a σ -finite positive measure on X and let $\mathcal{P}$ be a family of probability measures on X such that $P \ll \nu$ for all $P \in \mathcal{P}$. Then there exists $a_1, a_2, \dots \in [0, 1]$ with $\sum_i a_i = 1$ and $P_1, P_2, \dots \in \mathcal{P}$ such that

$$P \ll \sum_{i=1}^{\infty} a_i P_i \quad \text{for all } P \in \mathcal{P}.$$

Proof. By the previous Lemma, we only need to focus on the case that $\nu(X) = 1$. Let

$$\mathcal{Q} = \left\{ \sum_{i=1}^{\infty} a_i P_i \mid \sum_i a_i = 1, a_i \geq 0, P_i \in \mathcal{P} \right\}$$

be the set of countable convex combinations of measures in $\mathcal{P}$. Observe that $Q \ll \nu$ for all $Q \in \mathcal{Q}$. For $Q \in \mathcal{Q}$ let $f_Q = dQ/d\nu$ be the density of Q and let

$$S_Q = \{x \in X \mid f_Q(x) > 0\}.$$

We remark that f_Q and S_Q are almost surely well defined. Now let $c = \sup_{Q \in \mathcal{Q}} \nu(S_Q) \leq 1$. Let $Q_1, Q_2, \dots \in \mathcal{Q}$ be such that $\nu(S_{Q_i}) \rightarrow c$. Now there exists Q_0 such that

$$S_{Q_0} = \bigcup_{i=1}^{\infty} S_{Q_i}$$

and thus $\nu(S_{Q_0}) = c$. We claim that $P \ll Q_0$ for all $P \in \mathcal{P}$, thus completing the proof. Suppose it were not the case. Then there exists $A \subset X$ such that $P(A) > 0$ but $Q_0(A) = 0$. Now let

$$Q' = \frac{1}{2} (Q_0 + P).$$

We have that $Q_0 \ll Q'$ and so $S_{Q_0} \subset S_{Q'}$. Since $Q_0(A) = 0$ we have that $\nu(A \cap S_{Q_0}) = 0$. On the other hand

$$0 < P(A) = \int_A f_P d\nu \leq 2 \int_A f_{Q'} d\nu$$

and so this means that there is $A' \subset A$ with $\nu(A') > 0$ and $f_{Q'}(x) > 0$ for $x \in A'$. Thus

$$S_{Q_0} \cup A' \subset S_{Q'}$$

and so

$$\nu(S_{Q'}) \geq \nu(S_{Q_0}) + \nu(A') > c,$$

which is a contradiction. □

APPENDIX F. FUNCTIONS BETWEEN STANDARD BOREL SPACES

We start with a generalization of Proposition C.4 to maps into arbitrary standard Borel spaces rather than just $\mathbb{R}$. Recall that a standard Borel space is a measurable space $(X, \mathcal{B}_X)$ where X is a complete second countable metric space and $\mathcal{B}_X$ is the Borel σ -algebra on X .

Proposition F.1. Suppose that $(X, \mathcal{B}_X)$ and $(Y, \mathcal{B}_Y)$ are measurable space and let $(Z, \mathcal{B}_Z)$ be a standard Borel space. Let $S : X \rightarrow Y$ be measurable and let $\mathcal{B}_S = \{S^{-1}B \mid B \in \mathcal{B}_Y\}$ be the sub- σ -algebra of $\mathcal{B}_X$ induced by S . Suppose $T : X \rightarrow Z$ is $(\mathcal{B}_S, \mathcal{B}_Z)$ measurable. Then there exists a measurable $\psi : Y \rightarrow Z$ such that $T = \psi \circ S$.

Proof. We know from Corollary B.3 that there exists a homeomorphism $\iota : Z \rightarrow \tilde{Z}$ where $\tilde{Z}$ is a Borel subset of $[0, 1]^\mathbb{N}$. This means that we will now assume that $Z \subset [0, 1]^\mathbb{N}$. Let $\pi_i : [0, 1]^\mathbb{N} \rightarrow [0, 1]$ be the projection map and let $f_i = \pi_i \circ T : X \rightarrow [0, 1]$. Thus

$$T(x) = (f_1(x), f_2(x), \dots).$$

Now we can apply Proposition C.4 to $f_i : X \rightarrow [0, 1]$, which is $\mathcal{B}_S$ measurable as required, to get a measurable $\phi_i : Y \rightarrow [0, 1]$ such that $f_i = \phi_i \circ S$. Now let $\phi : Y \rightarrow [0, 1]^\mathbb{N}$ be given by

$$\phi(y) = (\phi_1(y), \phi_2(y), \dots).$$

Now let $\psi : Y \rightarrow Z$ be defined by

$$\psi(y) = \begin{cases} \phi(y), & \text{if } y \in \phi^{-1}(Z) \\ z_0, & \text{otherwise} \end{cases}$$

where $z_0 \in Z$ is fixed and arbitrary. Note that ψ is indeed measurable as $\phi^{-1}(Z)$ is measurable. Let us show that ψ now satisfies the desired property. If $x \in X$ then

$$\phi(S(x)) = (\phi_1(S(x)), \dots) = (f_1(S(x)), \dots) = T(x) \in Z$$

and thus $S(x) \in \phi^{-1}(Z)$. This means that $\psi(S(x)) = \phi(S(x)) = T(x)$, as desired. $\square$

Proposition F.2. Let $(X, \mathcal{B}_X)$, $(Y, \mathcal{B}_Y)$, $(Z, \mathcal{B}_Z)$ be standard Borel spaces and suppose that μ is a probability measure on $(X, \mathcal{B}_X)$. Suppose that $S : X \rightarrow Y$ and $T : X \rightarrow Z$ are Borel measurable and let $\mathcal{B}_S$ and $\mathcal{B}_T$ denote the sub- σ -algebras induced by these maps. Suppose that whenever $T(x) = T(x')$ we have that $S(x) = S(x')$. Then there T is μ -almost equal to a map that is $\mathcal{B}_S$ -measurable.

Proof. We only have to prove that if $f = \mathbb{1}_A$ for some $A \in \mathcal{B}_T$ then f is μ -almost surely equal to some map that is $\mathcal{B}_S$ -measurable. So fix such A . By assumption, we have that A is a union of atoms of $\mathcal{B}_S$, i.e., if we let $[x] = S^{-1}(S(x))$ denote the smallest element of $\mathcal{B}_S$ that contains x (see ???) then

$$A = \bigcup_{x \in A} [x].$$

Now we have (by ???) that there exists $X' \in \mathcal{B}_S$ with $\mu(X') = 1$ such that if $\mu_x = \mu_x^{\mathcal{B}_S}$ denotes the conditional probability then $\mu_x^{\mathcal{B}_S}([x]) = 1$ for each $x \in X'$. We have (by WHAT??) that the map $\tilde{f} : X' \rightarrow \mathbb{R}$ given by

$$\tilde{f}(x) = \int_X \mathbb{1}_A(u) d\mu_x(u)$$

is $\mathcal{B}_S$ -measurable. Now observe that for $x \in X'$ we have that if $x \in A$ then $[x] \subset \mathcal{A}$ and so this integral is 1. Otherwise, if $x \in X'$ and $x \notin A$ then $[x] \cap \mathcal{A} = \emptyset$ and so this integral is 0. Thus we have shown that for $x \in X'$ we have that $f(x) = \tilde{f}(x)$ and $\tilde{f}$ is $\mathcal{B}_S$ measurable on $X' \in \mathcal{B}_S$, which completes the proof. $\square$